

Core Products

Three feature concepts, feasibility and MVP economics

Product strategy

Early-stage CTO review

5 September 2026

1 Executive decision

The product should not begin as an all-in-one fitness platform. The reviewed documents consistently support a narrower wedge: **help a self-coached recreational lifter make one explainable decision before a planned strength session**. The strongest feature is therefore the **Evidence-Backed Session Decision**, supported by a deliberately thin training ledger and manual or simulated recovery inputs.

Recommendation: Build Feature 2 as the Challenge centrepiece. Treat Feature 1 as enabling infrastructure inside it, not as a separate first product. Defer Feature 3 until athletes repeatedly use and trust the individual decision loop.

2 Assessment method and cost basis

Technical feasibility is judged against a pilot-quality vertical slice, not a production platform with every proposed connector. **High** means mature patterns and accessible inputs can deliver the user value with limited uncertainty. **Medium** means the feature is feasible but depends on data quality, safety rules, permissions or a long-lived workflow that requires meaningful validation.

Cost ranges are planning allowances, not supplier quotes. They assume one senior TypeScript full-stack contractor at €650–€850/day, fractional product/design at €500–€700/day, fractional QA/security at €550–€750/day and 15% contingency. They exclude VAT, founder time, legal/DPIA work, clinical validation, provider partnership fees, native HealthKit/Health Connect work and post-pilot operations.

3 Feature 1 — Smart Training Ledger + Ghost Sets

Dimension	Assessment
User job	Turn familiar workout notation into a structured block/session record and show the last comparable set beside today's editable fields.
Minimum scope	Block → session → movement → set; paste parser with ambiguity confirmation; prior comparable-session retrieval; ghost sets; append-only completed session.
Technical feasibility	HIGH Mature relational and UI patterns. The main unknown is parser usability, not core engineering viability.
Effort	4–5 calendar weeks; approximately 26–34 person-days.
Rough build cost	€18k–€32k under the stated assumptions.

Dimension	Assessment
Strategic upside	Creates the canonical training context required by later recommendations and has immediate utility without wearable access.
Primary weakness	Workout logging and progression history are already served by Hevy, Strong and others. Alone, this is less differentiated and less visibly agentic.

This is the lowest-risk build and should be present inside any prototype. It is not the strongest standalone launch proposition because its clearest benefits—faster logging and visible progression—are familiar category features. The unique value appears only when this record is used to justify a decision.

4 Feature 2 — Evidence-Backed Session Decision

Dimension	Assessment
User job	Before a planned session, decide whether to progress, repeat, reduce volume or defer a hard top set—and understand why.
Minimum scope	Active block and comparable-session context; manual/CSV recovery summary; data-completeness check; deterministic policy; one strict-schema proposal; evidence, confidence, counterfactual and gaps; approve/modify/keep-original; prepared session and linked outcome.
Technical feasibility	MEDIUM Feasible with managed services, but explanation quality, safety boundaries, partial data and outcome capture must be tested.
Effort	6–8 calendar weeks; approximately 41–56 person-days.
Rough build cost	€29k–€52k under the stated assumptions.
Strategic upside	Owns the precise problem identified across the source pack and demonstrates observe → verify → propose → approve → act → remember in one flow.
Primary weakness	The recommendation can appear authoritative despite imperfect wearable evidence. The product needs bounded actions, non-clinical language, visible uncertainty and a deterministic fallback.

This feature best combines problem fit, differentiation and Challenge impact. Its feasibility improves materially when live wearable integrations are removed from the critical path. Synthetic/de-identified scenarios, manual input and a small CSV import are acceptable prototype inputs under the Challenge boundary, which explicitly permits responsibly simulated data and mock integrations. [\[1\]](#)

5 Feature 3 — Coach Review Queue + Decision History

Dimension	Assessment
User job	Let a coach review flagged proposals using only the evidence the athlete has authorised, approve or amend the plan and preserve decision authority.

Dimension	Assessment
Minimum scope	Coach/athlete tenancy; explicit access grant; review queue; minimal evidence view; approve/change/reject; rationale; notifications; immutable decision and access audit.
Technical feasibility	MEDIUM Standard software patterns, but multi-tenant authorisation and sensitive-data minimisation raise the release bar.
Effort	7–9 calendar weeks; approximately 44–63 person-days.
Rough build cost	€31k–€59k under the stated assumptions.
Strategic upside	Improves the buyer signal, creates a B2B2C path and turns low-confidence recommendations into accountable human review.
Primary weakness	It depends on athlete activity, coach recruitment and a trusted decision loop. Building it first risks an empty queue and premature enterprise complexity.

This is the best second wedge after individual demand is demonstrated. The authorisation model is not optional: Supabase’s current guidance emphasises enabling and testing Row Level Security on exposed tables, with explicit grants and policies for each operation. [\[2\]](#)

6 Comparison matrix

Feature	Feasibility	Effort	Cost	Score	Decision
Smart Training Ledger	High	4–5 wks	€18k–€32k	7.0/10	Enabler
Evidence-Backed Session Decision	Medium	6–8 wks	€29k–€52k	8.4/10	Select
Coach Review + History	Medium	7–9 wks	€31k–€59k	7.1/10	Phase 2

The score weights problem fit 30%, visible agentic value 20%, buildability 20%, differentiation 15%, data/trust risk 10% and commercial leverage 5%. Scores are strategic planning judgements, not market measurements.

Weighted criterion	Ledger	Decision	Coach
Problem fit — 30%	8	10	7
Agentic demo value — 20%	4	10	8
Buildability — 20%	9	6	6
Differentiation — 15%	5	8	7
Data/trust risk — 10%	9	6	7
Commercial leverage — 5%	6	8	9

7 Product sequencing

Stage	Ship	Gate to continue
Now	Evidence-Backed Session Decision with minimum ledger	Five usability tests can explain the recommendation without help; ten golden scenarios pass policy and evidence checks; the demo runs without provider APIs.
Next	Harden ledger and outcome history	Design partners repeatedly prepare sessions and record outcomes; time-to-decision and trust are instrumented.
Later	Coach review workspace	At least two coaches and a live athlete cohort demonstrate a recurring review job and willingness to pay.

8 Sources and internal basis

[1] [TechIreland, National AI Challenge](#). The supplied 2026 curated challenge statements were used for prototype and human-accountability criteria.

[2] [Supabase, “Row Level Security”](#), accessed 5 September 2026.

The product assessment also reviewed every document supplied with the brief, including the pitch deck, idea validation, competitive assessment, methodology/ranked list, problem statement, value proposition, USP, elevator pitch, speech, recommendations, prior development plan and agentic workflow specification.