

Competitive Assessment — Training Ledger vs TechIreland National AI Challenge Field

Scoring Methodology

Each entry is scored on seven dimensions aligned with the official judging criteria and the 2025 keynote philosophy. Each dimension is scored 1-10.

Dimension	Weight	What judges look for
Agentic Design	20%	Multi-step workflow, memory, escalation, human-in-the-loop
Technical Depth	15%	Architecture quality, production-readiness, not just a wrapper
Problem Clarity	15%	Named user, quantified cost, root cause not symptom
Commercial Viability	15%	Buyer, pricing, route to revenue
Responsible AI	10%	Governance, consent, explainability, safety boundaries
Demo Impact	15%	Wow moment, end-to-end workflow visible in 3 minutes
Domain Expertise	10%	Do what you know — founder lives the problem

Training Ledger Score

Dimension	Score	Justification
Agentic Design	9/10	9-agent pipeline with typed handoffs, durable workflows, approval gates, deterministic policy before and after model. Textbook agentic architecture.
Technical Depth	8/10	Full production blueprint (serverless, TypeScript, Supabase, Inngest, OpenAI SDK). Not a thin wrapper — only 1 of 9 agents uses an LLM. Architecture diagrams, state machine, tool contracts all specified.
Problem Clarity	8/10	Clear user (serious lifters), quantified (1.97M Irish sport participants, 538M wearables), root cause identified (data silos, missing training context). Slightly weaker: no single enterprise buyer.
Commercial Viability	7/10	99 euro per year subscription with clear positioning. Coach tier adds B2B. But B2C is harder to prove in 2 weeks than B2B enterprise. No LOI or pilot commitment yet.
Responsible AI	9/10	Consent-first design, GDPR Article 9 awareness, non-clinical boundary, policy-before-model, prohibited claims registry, data minimisation, audit trail. Exceptional for a hackathon entry.
Demo Impact	7/10	Notes-to-structured-session parser plus ghost sets plus evidence card is compelling. But no live wearable data flowing in demo (API keys needed). Mock data demo is strong but not visceral.
Domain Expertise	10/10	Founder IS the user. Tracks blocks, wears devices, logs food, does the manual reconciliation daily. This is the do what you know unfair advantage the 2025 keynote described.

Weighted Total: 8.3 / 10

Comparison with Previous Winners

Entry	Year	Agentic	Technical	Problem	Commercial	Responsible AI	Demo	Domain	Weighted
Training Ledger	2026	9	8	8	7	9	7	10	8.3
Bias Busters	2024	4	6	7	4	10	6	7	6.0
HandSpeak AI	2025	5	7	8	5	7	9	7	6.7
LaunchLoop	2025	7	7	7	7	6	7	6	6.8
LibertyLumens	2025	6	6	7	6	7	7	6	6.4

Context Notes

2024 was pre-agents era. Bias Busters won with ethical AI focus — no agentic workflow expected. The bar has risen dramatically.

2025 introduced agentic thinking but most entries were still single-model wrappers with basic tool use. HandSpeak AI won partly on demo impact (sign language is visually compelling).

2026 explicitly demands agents. The curated challenges all specify multi-step workflows with human handoffs. Training Ledger's 9-agent architecture with typed contracts and durable orchestration is significantly more sophisticated than anything shown in previous years.

Wildcard risk: Judges may unconsciously favour entries that respond to a curated challenge (clearer buyer signal). Training Ledger must emphasise: the user IS the buyer, the market is quantified, and the agentic design is deeper than most curated-challenge responses will achieve.

Comparison with Likely 2026 Curated-Challenge Entries

The strongest curated challenges for 2026 (based on clarity, buyer, and agentic fit) are likely:

- Challenge 1 (Late-Payment Resolution Agent) — clear SME buyer, simple chase workflow
- Challenge 9 (Compliant Rota Agent) — hospitality, scheduling is well-understood
- Challenge 10 (Tender Qualification Agent) — public procurement, high value
- Challenge 13 (Home Retrofit Sequencing) — climate and housing, government interest

Factor	Curated entries (typical)	Training Ledger (wildcard)
Named buyer	Enterprise or government (strong)	Consumer plus coach (weaker signal)
Agentic depth	Usually 3-5 steps, basic tool use	9 agents, typed contracts, durable orchestration
Domain expertise	Often assembled team, may lack deep domain	Founder lives the problem daily
Demo viscosity	Synthetic data, mock integrations	Same constraint, but personal data makes it real
Responsible AI	Often an afterthought	Core architectural principle
Post-challenge path	Pilot customer conversation	Subscription product with clear unit economics

Overall Assessment

Training Ledger scores in the top tier of what this competition has ever seen on agentic design, responsible AI, and domain expertise. Its primary weakness is demo viscosity (no live wearable data in a 2-week sprint) and the B2C buyer signal (judges may prefer enterprise). If the demo lands — particularly the notes-parser to ghost-sets to evidence-card flow — this is a strong national finalist and a credible winner.

Probability of reaching national final: 75-85% (assuming competent presentation)

Probability of winning overall: 25-35% (dependent on demo execution and judge composition)

The main competitive threat is a well-executed insurance or finance challenge entry with a named enterprise sponsor providing a letter of intent. Training Ledger's counter: deeper agentic architecture, stronger responsible AI story, and authentic domain expertise that no assembled team can fake.